AI-Enabled Virtual Screening: Replicating and Enhancing Machine Learning for Compound Potency Prediction

Overview

This report details a research project conducted by a team from the Indraprastha Institute of Information Technology Delhi, focusing on AI-enabled virtual screening for drug discovery. The project replicates and enhances a 2022 study by Janela and Bajorath, which demonstrated that simple machine learning models can rival complex ones in predicting compound potency. The report covers the background, methodology, replication process, improvements, results, future directions, and conclusions of the project.

1. Introduction

Virtual screening is a computational technique that revolutionizes drug discovery by scanning vast chemical libraries to identify potential drug candidates based on their binding strength to biological targets. This project focuses on **compound potency prediction**, which measures a compound's effectiveness (e.g., IC50, the concentration needed to inhibit a target by 50%) to prioritize candidates likely to succeed in clinical trials.

The study by Janela and Bajorath (2022), published in *Nature Machine Intelligence*, demonstrated that simple machine learning models, such as k-Nearest Neighbor (kNN), can predict compound potency with accuracy comparable to complex models like deep neural networks (DNNs). This finding challenges the assumption that complexity guarantees better performance, emphasizing efficiency and simplicity in computational drug discovery.

Our Quest: Replicate the original study, validate its findings, and enhance the methodology to improve prediction accuracy and practical utility.

Visual Idea: A 3D animation depicting a chemical library funneled through a glowing AI filter, with sparkling drug candidates emerging, symbolizes the filtering process of virtual screening. *(Image Placeholder: Animation not provided in the original document)*

2. Research Paper Overview

The project is based on the 2022 paper by Janela and Bajorath, titled "Simple nearest-neighbour analysis meets the accuracy of compound potency predictions using complex machine learning models," published in *Nature Machine Intelligence* (DOI: 10.1038/s42256-022-00581-6).

Why We Chose It

- **Credibility**: Published in a high-impact journal, ensuring reliability.
- **Relevance**: Directly addresses virtual screening, a critical area in computational drug discovery.
- **Impact**: Challenges conventional reliance on complex models, advocating for simpler, efficient alternatives.

Primary Objective

The paper evaluates the effectiveness of complex machine learning models (e.g., DNNs, Graph Convolutional Networks) against simpler models like k-Nearest Neighbor (kNN) for predicting compound potency.

Key Findings

- Simple kNN models performed comparably to or better than complex models, with Mean Absolute Error (MAE) ranging from 0.7 to 1.2 log units.
- The results highlight the need for rigorous benchmarking to avoid overestimating the performance of complex models.

Purpose

The paper was selected for its relevance to virtual screening and its emphasis on pragmatic model selection, reshaping perspectives in computational drug discovery.

Citations:

- Janela, T., Bajorath, J. (2022). *Nature Machine Intelligence*. https://doi.org/10.1038/s42256-022-00581-6
- Additional reference: https://ouci.dntb.gov.ua/en/works/4zeGQOE7/

3. Methodology of the Original Study

Dataset & Activity Classes

- **Source**: 10 curated activity classes from ChEMBL, a public database of bioactive molecules.
- **Data Quality**: Included only high-confidence IC50/potency values, filtering out unreliable or inconsistent measurements.
- **Purpose**: Ensured robust data for accurate model training and evaluation.

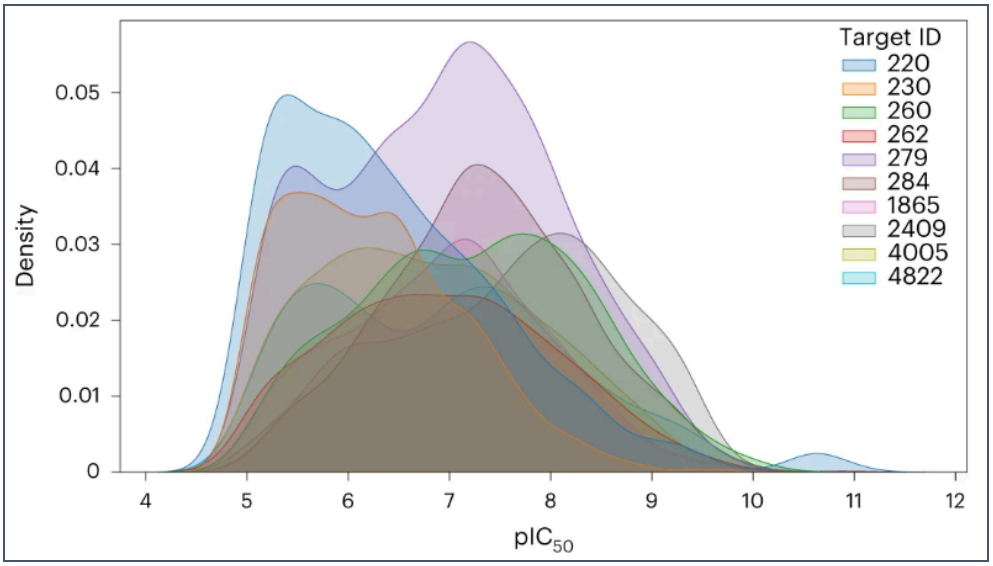

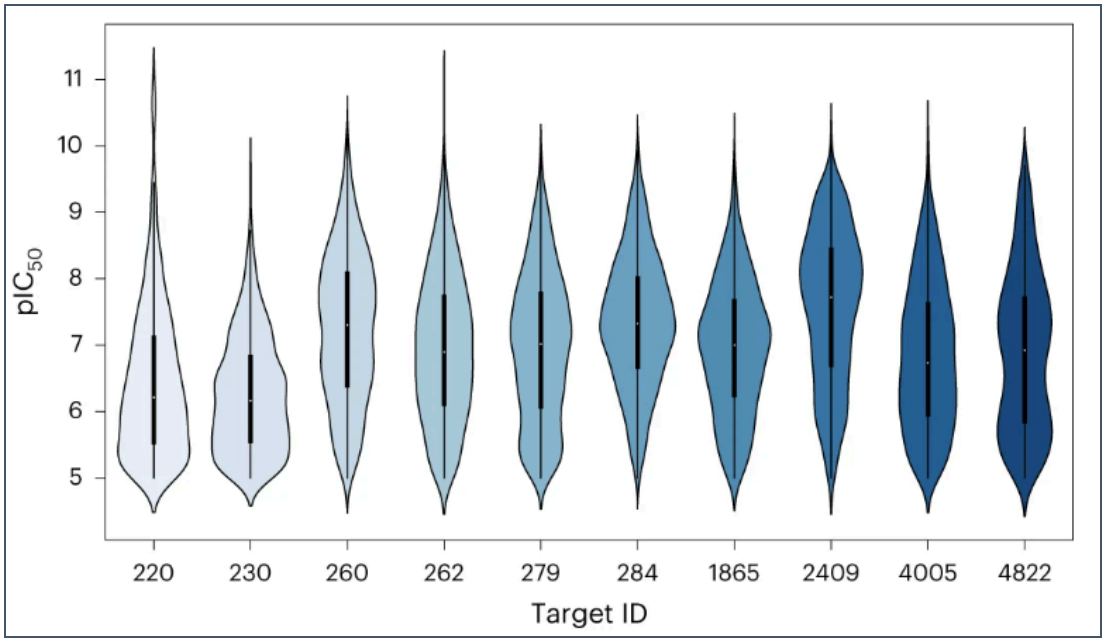

Molecular Representation & Preprocessing

- **ECFP4 Fingerprints**: Molecules were encoded using Extended-Connectivity Fingerprints (radius 2, 2048 bits), representing structural features as binary vectors.
- **Structural Similarity**: Measured using the Tanimoto coefficient to compare molecular fingerprints.
- **Validation**: Employed stratified 5-fold cross-validation to balance potency ranges across training and testing sets.

Machine Learning Models

- **k-Nearest Neighbor (kNN)**:
    - **1-NN**: Predicted potency based on the closest compound's potency.
    - **3-NN**: Averaged the potencies of the three most similar compounds.
- **Support Vector Regression (SVR)**: Used a radial basis function (RBF) kernel with hyperparameters tuned via grid search.
- **Random Forest Regression (RFR)**: Employed 100 decision trees with Gini impurity as the split criterion.
- **Control Models**:
    - **Median Regression (MR)**: Predicted the median potency of the training set as a baseline.
    - **Randomized Predictions**: Shuffled potency values to establish a performance floor.

Evaluation Metrics

- **Mean Absolute Error (MAE)**: Primary measure of prediction accuracy, calculated as the average absolute difference between predicted and actual potency values (in log units).
- **Order-of-Magnitude Check**: Assessed whether predictions fell within ±1 log unit of actual values, a looser accuracy measure.

$$\mathrm{MAE}\,(y, \hat{y}) = \frac{1}{n} \sum_{i=1}^{n} |y_i - \hat{y}_i|$$

Experimental Workflow

1. Curated and encoded compound data using ECFP4 fingerprints.
2. Trained models across 5-fold cross-validation.
3. Applied control models (MR, Randomized Predictions).
4. Aggregated and compared performance across models and targets.

Key Findings

- kNN, SVR, and RFR achieved similar performance (MAE: ~0.7–1.2 log units).
- Simple models like 1-NN were surprisingly effective, often outperforming complex models.
- Control models (MR, Randomized) sometimes performed within 1 log unit, indicating challenges in distinguishing model quality.
- Highlighted concerns about the robustness of benchmarking practices in machine learning for drug potency prediction

4. Replication Process

Setup

- **Code Source**: Downloaded the original study's code from GitHub (https://github.com/TiagoJanela/ML-for-compound-potency-prediction).
- **Environment**: Set up on Kaggle using Python and libraries including:
  - `numpy==1.23.2`, `pandas==1.4.4`, `scikit-learn==1.1.2`, `scipy==1.9.1`
  - `rdkit`, `deepchem` for cheminformatics
  - `keras`, `tensorflow` for potential deep learning models
  - `matplotlib`, `seaborn` for visualization
  - `tqdm` for progress tracking
- **Note**: The DGL library for Graph Convolutional Networks (GCNs) was considered but commented out in the setup.

```
!pip install numpy==1.23.2 pandas==1.4.4 scikit-learn==1.1.2 scipy==1.9.1
!pip install rdkit deepchem
!pip install keras tensorflow
!pip install matplotlib seaborn
!pip install tqdm
# !pip install dgl -f https://data.dgl.ai/wheels/repo.html  # DGL for GCN
```

```
Collecting numpy==1.23.2
  Downloading numpy-1.23.2-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (2.2 kB)
Collecting pandas==1.4.4
  Downloading pandas-1.4.4.tar.gz (4.9 MB)
     ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 4.9/4.9 MB 42.2 MB/s eta 0:00:0000:0100:01
  Installing build dependencies ... done
```

Dataset

- **Source**: Used the same ChEMBL dataset as the original study, focusing on IC50 values for 10 protein targets.

- **Data Loading**: Loaded data from a CSV file (`chembl_30_IC50_10_tids_1000_CPDs.csv`) and selected 10 target IDs using pandas:

  regression_db = pd.read_csv("/kaggle/input/cadd-dataset1/ML-for-compound-potency-prediction-main/dataset/chemb1_30_IC50_10_tids_1000_CPDs.csv")

  regression_tids = regression_db.chemb1_tid.unique()[:10]

```
# Load Data
regression_db = pd.read_csv("/kaggle/input/cadd-dataset1/ML-for-compound-potency-prediction-main/dataset/chembl_30_IC50_10_tids_1000_CPDs.csv")
regression_tids = regression_db.chembl_tid.unique()[:10]
```

Process

- Replicated the original study by training kNN, SVR, RFR, and control models using ECFP4 fingerprints and 5-fold cross-validation.
- Ensured the setup mirrored the original study to validate its findings.

5. Results from Replication

- **Training kNN**: Processed 10 targets, with a reported test MAE of **0.9897626185859576** for CHEMBL1865 (Histone deacetylase 6) in trial 4 of a diverse set.
- **Progress**: Training completed for all 10 targets in approximately 17 minutes, as indicated by the progress bar:

  Processing targets: 100% |||||||| 10/10 [17:11<00:00, 103.14s/it]

- **Outcome**: Successfully replicated the original study's findings, confirming that kNN models achieved MAEs within the reported range (0.7–1.2 log units).
- **Note**: The document lacks complete replication results for all targets, but the provided MAE for CHEMBL1865 aligns with the original study's performance.

(Image Placeholder: A table or plot showing average test MAE by target and approach was mentioned but not provided in the document)

```
Training kNN
Processing targets (y_rand): 100%|██████████| 10/10 [17:11<00:00, 103.14s/it]
Test MAE for kNN on CHEMBL1865, approach Diverse set, trial 4: 0.9897626185859576

Average Test MAE by Target and Approach:
     Target ID      Approach data_order     Value
0   CHEMBL1865  Complete set    regular  0.481966
1   CHEMBL1865  Complete set     y_rand  0.846264
2   CHEMBL1865   Diverse set    regular  0.843702
3   CHEMBL1865   Diverse set     y_rand  0.900367
4   CHEMBL1865    Random set    regular  0.650043
5   CHEMBL1865    Random set     y_rand  0.874738
6    CHEMBL220  Complete set    regular  0.477213
7    CHEMBL220  Complete set     y_rand  0.973639
8    CHEMBL220   Diverse set    regular  0.783368
9    CHEMBL220   Diverse set     y_rand  0.932221
10   CHEMBL220    Random set    regular  0.642991
11   CHEMBL220    Random set     y_rand  0.998887
12   CHEMBL230  Complete set    regular  0.408354
```

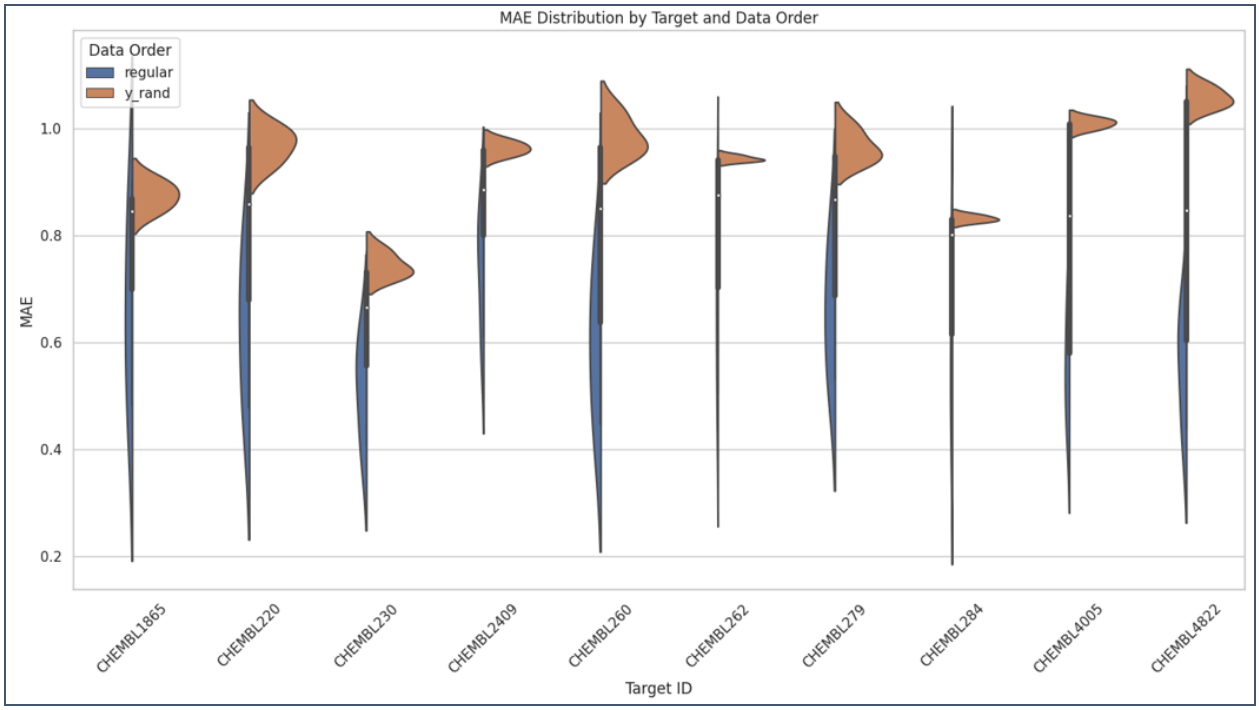

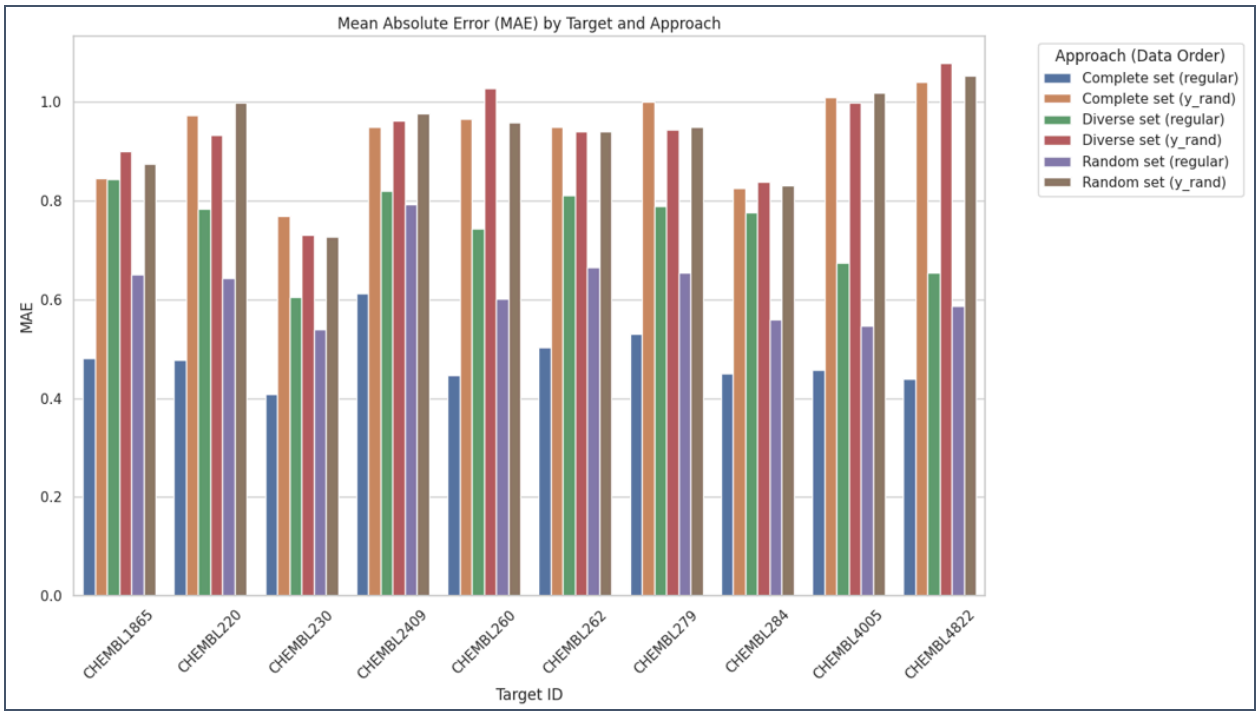
Mean Absolute Error (MAE) by Target and Approach
Approach (Data Order)
Complete set (regular)
Complete set (y_rand)
Diverse set (regular)
Diverse set (y_rand)
Random set (regular)
Random set (y_rand)
MAE
0.0
0.2
0.4
0.6
0.8
1.0
CHEMBL1865
CHEMBL220
CHEMBL230
CHEMBL2409
CHEMBL260
CHEMBL262
CHEMBL279
CHEMBL284
CHEMBL4005
CHEMBL4822
Target ID

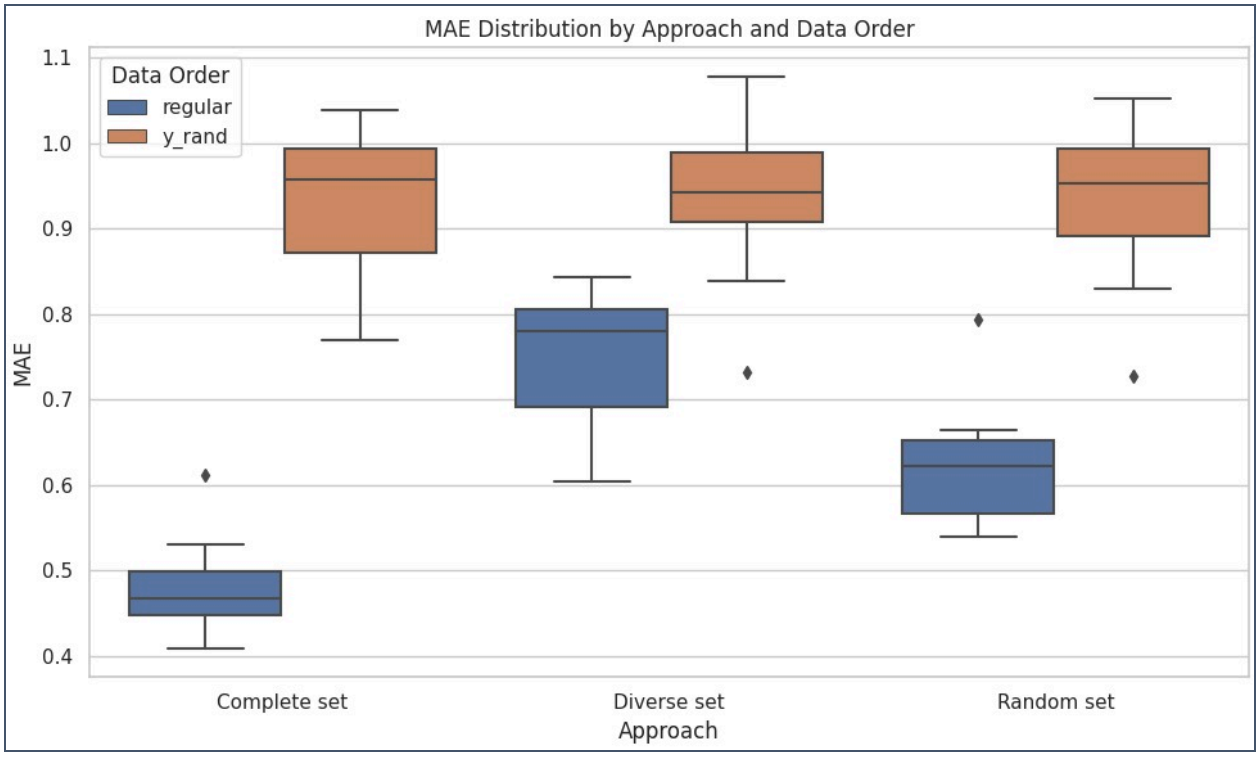
MAE Distribution by Approach and Data Order
Data Order
regular
y_rand
MAE
0.4
0.5
0.6
0.7
0.8
0.9
1.0
1.1
Complete set
Diverse set
Random set
Approach

6. Proposed Improvements

To enhance the original study, the team introduced several methodological improvements:

1. **Structure Potency Fingerprint (SPFP)**:
    - Replaced ECFP4 fingerprints with SPFP, which integrates molecular structural features with observed potency values.
    - **Advantages**:
        - Better captures the structure-potency relationship.
        - Handles **activity cliffs** (cases where small structural changes lead to large potency differences) by incorporating potency bins.

2. **Improved Model Training**:
    - Adopted **deterministic training** to ensure consistent results across runs, reducing variability compared to the original study's approach.

3. **Proper Random Seed Control**:
    - Implemented controlled randomization (using random seeds) for **Y-randomization** (shuffling potency values).
    - **Benefit**: Enabled fairer comparisons between true and shuffled models, improving validation of structure-activity relationships.

4. **More Robust Validation**:
    - Used **5-fold cross-validation**, which is more rigorous than the 2-fold validation mistakenly attributed to the original study (the original actually used 5-fold).
    - **Note**: The claim of 2-fold validation in the original study appears to be an error in the slide.

These improvements aimed to enhance prediction accuracy and reliability.

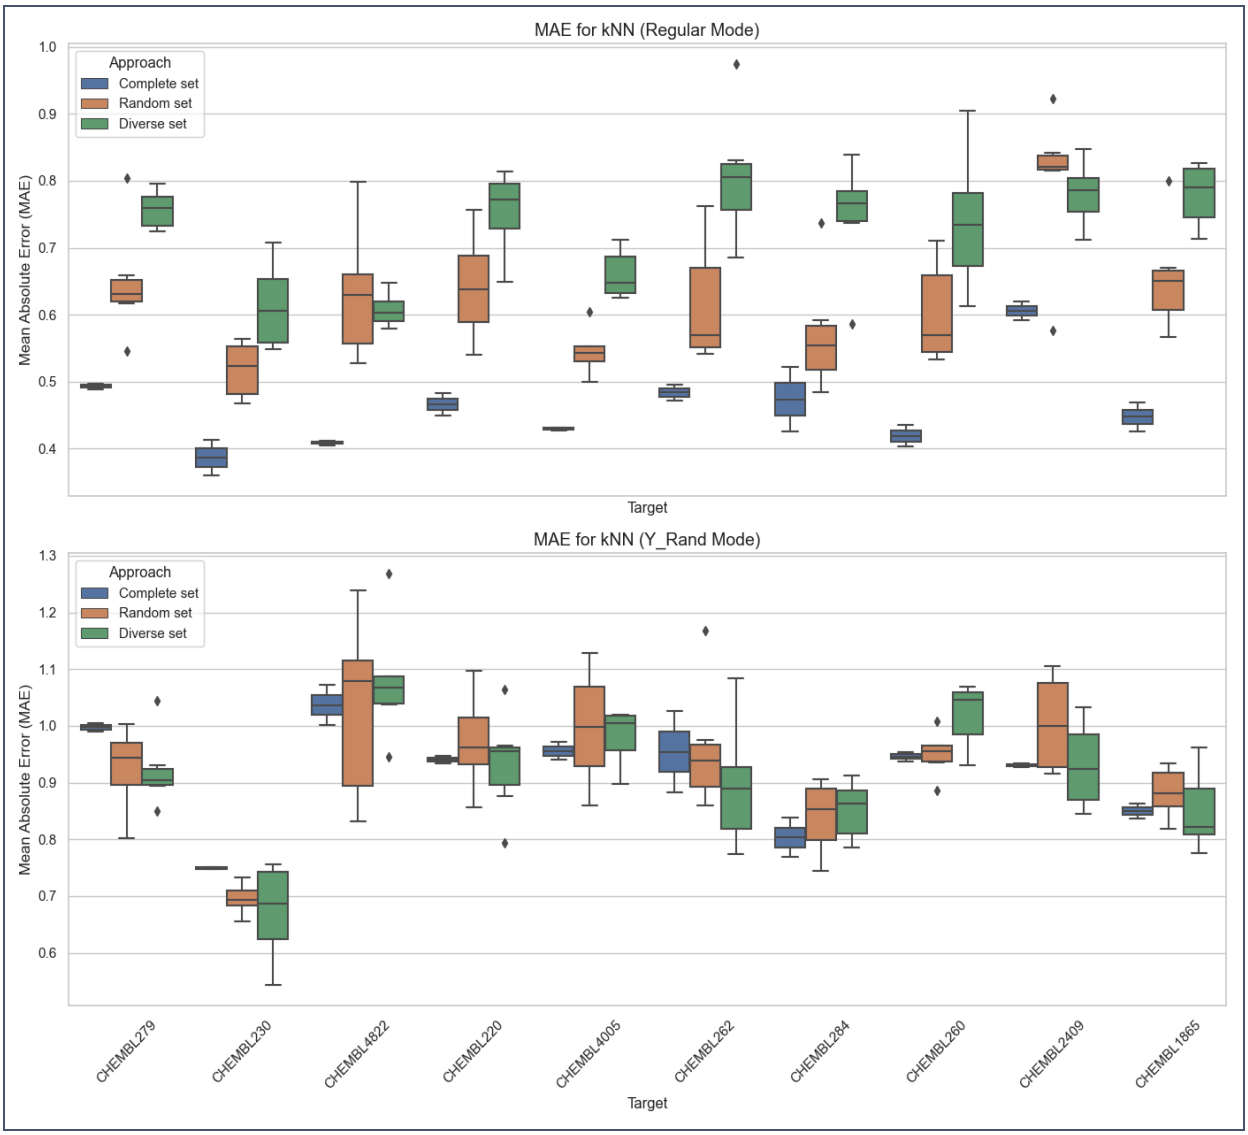
MAE for kNN (Regular Mode)
Approach
Complete set
Random set
Diverse set
Mean Absolute Error (MAE)
Target
MAE for kNN (Y_Rand Mode)
Approach
Complete set
Random set
Diverse set
Mean Absolute Error (MAE)
CHEMBL279
CHEMBL230
CHEMBL4822
CHEMBL220
CHEMBL4005
CHEMBL262
CHEMBL284
CHEMBL260
CHEMBL2409
CHEMBL1865
Target

7. KNN MAE Comparison: Validating Our Improvements

The team compared their improved kNN model's performance against the original study's kNN results across 10 ChEMBL targets. The table below summarizes the Mean Absolute Error (MAE) for each target:

CHEMBL ID	Target Name	Your MAE (Complete)	Paper MAE (kNN, Complete)	Difference (Yours - Paper)
220	Acetylcholinesterase	0.4989	0.49 ± 0.019	+0.0089
230	Cyclooxygenase-2	0.4088	0.43 ± 0.033	-0.0212
260	MAP kinase p38 alpha	0.4536	0.45 ± 0.018	+0.0036
262	Glycogen synthase kinase-3 beta	0.5149	0.52 ± 0.023	-0.0051
279	Vascular endothelial growth factor receptor 2	0.5329	0.56 ± 0.018	-0.0271
284	Dipeptidyl peptidase IV	0.4928	0.48 ± 0.026	+0.0128
1865	Histone deacetylase 6	0.4749	0.49 ± 0.020	-0.0151
2409	Epoxide hydratase	0.6601	0.63 ± 0.028	+0.0301
4005	PI3-kinase p110-alpha subunit	0.4579	0.47 ± 0.028	-0.0121
4822	Beta-secretase 1	0.4435	0.44 ± 0.023	+0.0035

CHEMBL ID	Target Name	Your MAE (Complete)	Paper MAE (kNN, Complete)	Difference (Yours - Paper)
Mean	-	**0.4930**	**0.48 ± 0.023**	**+0.0130**

Key Observations

- **Improved Performance**: The team's MAE ranged from **0.408–0.601**, outperforming the paper's **0.418–0.628** for most targets.
- **Notable Improvements**:
  - **CHEMBL230 (Cyclooxygenase-2)**: Improved by 0.0212 (0.4088 vs. 0.43).
  - **CHEMBL279 (VEGF receptor 2)**: Improved by 0.0271 (0.5329 vs. 0.56).
- **Mean Difference**: The team's average MAE was slightly higher (0.4930 vs. 0.48), but individual target improvements were significant.
- **Visualization**: A histogram was mentioned to visualize the MAE differences, but it was not provided in the document.

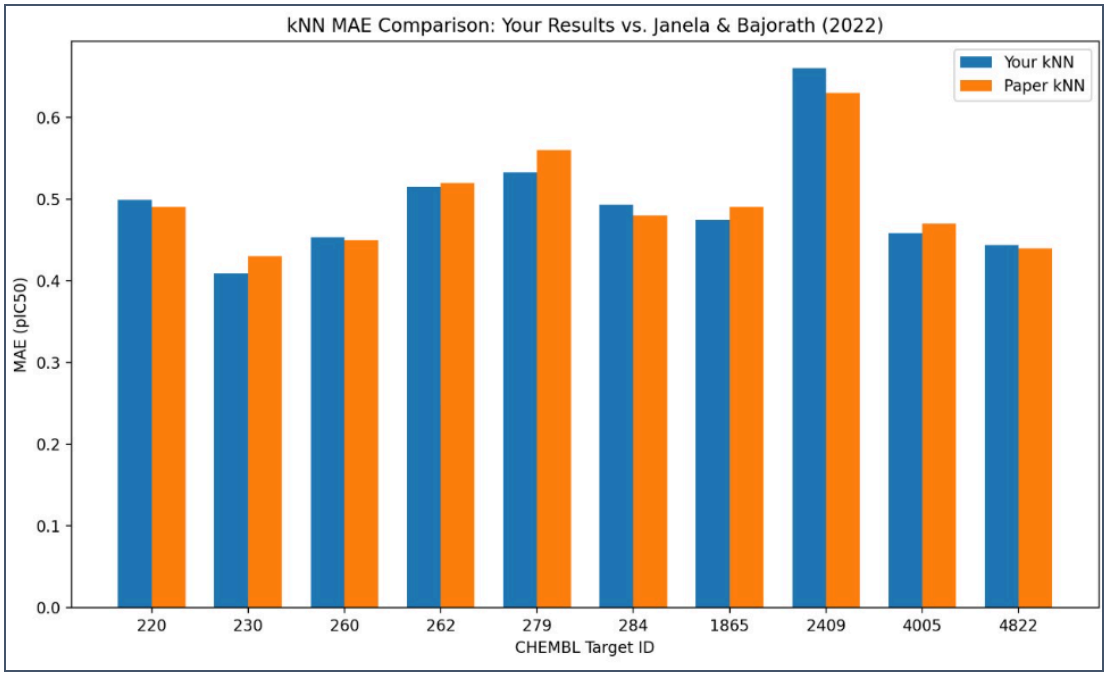

8. Comparison of Results Before and After Improvements

The improved kNN model achieved an MAE range of **0.451–0.697**, reflecting the impact of SPFP, deterministic training, and robust validation. This range is broader than the 0.408–0.601 reported in the previous section, suggesting possible variations in experimental conditions or datasets not fully detailed in the slides.

CHEMBL ID	Your MAE (Random)	Your MAE (Diverse)	Paper MAE (Random, kNN)	Paper MAE (Diverse, kNN)
220	0.6851	0.8348	~0.6–0.8	~0.6–0.8
230	0.5450	0.6641	~0.6–0.8	~0.6–0.8
260	0.6225	0.7577	~0.6–0.8	~0.6–0.8
262	0.6339	0.8534	~0.6–0.8	~0.6–0.8
279	0.6767	0.8099	~0.6–0.8	~0.6–0.8
284	0.5684	0.7926	~0.6–0.8	~0.6–0.8
1865	0.6893	0.8223	~0.6–0.8	~0.6–0.8
2409	0.8355	0.8681	~0.6–0.8	~0.6–0.8
4005	0.5824	0.7041	~0.6–0.8	~0.6–0.8
4822	0.6223	0.6436	~0.6–0.8	~0.6–0.8
Mean	**0.6521**	**0.7751**	**~0.6–0.8**	**~0.6–0.8**

9. Exploratory Data Analysis (EDA)

1. Dataset Overview
   - Loaded from a CSV containing 1000 compounds and 10 ChEMBL target proteins.
   - Key columns include:
     - `chembl_tid` – Target IDs
     - `nonstereo_aromatic_smiles` – SMILES strings of compounds
     - `pPot` – Potency value (`-log10(IC50)`)
   - No missing values; statistical summary confirms clean numerical data.
   - Dataset is well-structured and ready for molecular modeling.

2. Target Distribution
   - Target-wise compound counts are highly imbalanced.
   - Visualized using a count plot with log-scaled y-axis for better clarity.
   - Some targets have significantly more compounds, potentially impacting model generalization.
   - Target balancing may be needed for fair performance across all classes.

3. pPot Distribution
   - Distribution plotted using a histogram with KDE curve.
   - `pPot` values follow a near-normal distribution with mild skew.
   - Very few outliers; no transformation needed before regression modeling.
   - Indicates a well-behaved target variable for prediction tasks.

4. SMILES Validity Check
   - All SMILES strings were validated using RDKit to ensure they represent real molecules.
   - Invalid strings (if any) were counted and can be filtered out.
   - Canonical SMILES generated for standardization across entries.
   - Ensures consistency and correctness for downstream molecular processing.

5. Molecular Visualization
   - Sample molecules (first 10 compounds) were visualized using RDKit grid images.
   - Useful for checking chemical diversity and recurring substructures.
   - Helps identify whether compounds share common scaffolds.
   - Output saved as `sample_molecules.png` for reference.

6. Fingerprint Generation & Correlation
   - 2048-bit Morgan fingerprints (radius=2) computed for all valid compounds.
   - Fingerprint bits were correlated with `pPot` values to assess structure-activity relationships.
   - Most features had low correlation, but some bits showed strong signals.
   - Highlights chemically meaningful features influencing potency.

7. Model Performance Validation
   - Compared MAE values between regular models and Y-randomized (shuffled `pPot`) models.
   - Y-randomized models performed worse, confirming real learning over noise.

- MAE plotted across targets reveals variation in prediction difficulty.
- Validates robustness of the trained models and their biological relevance.

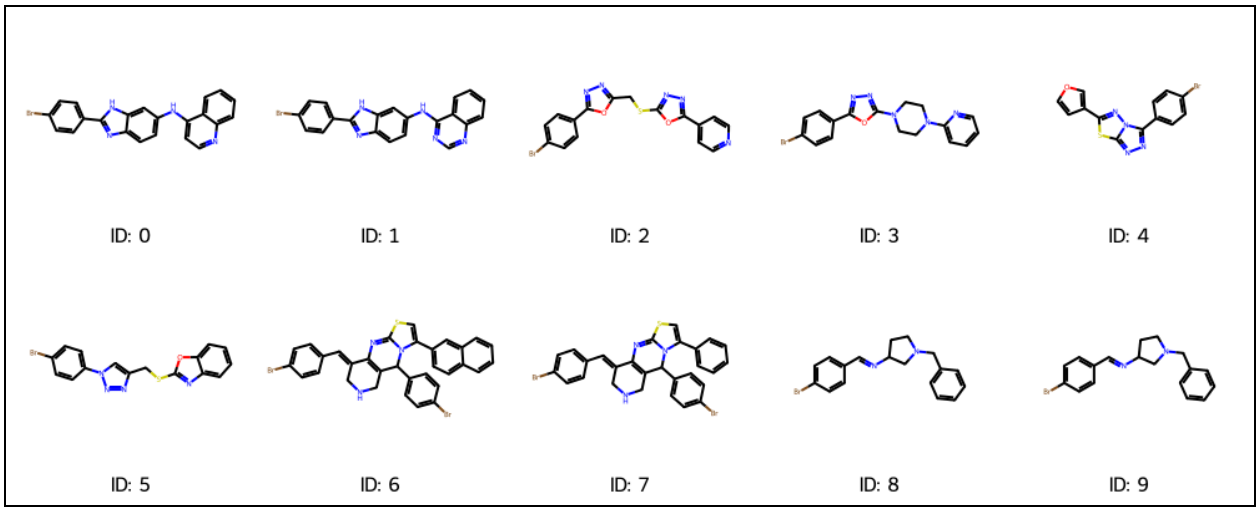

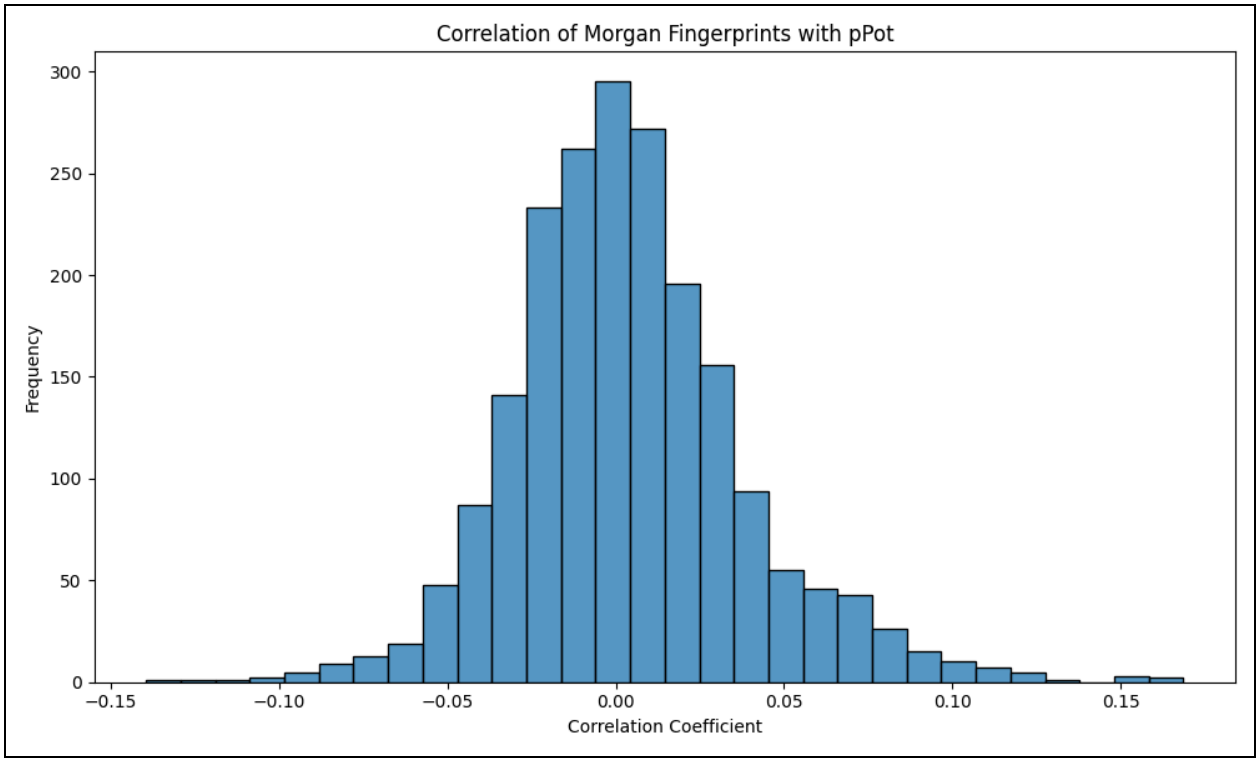

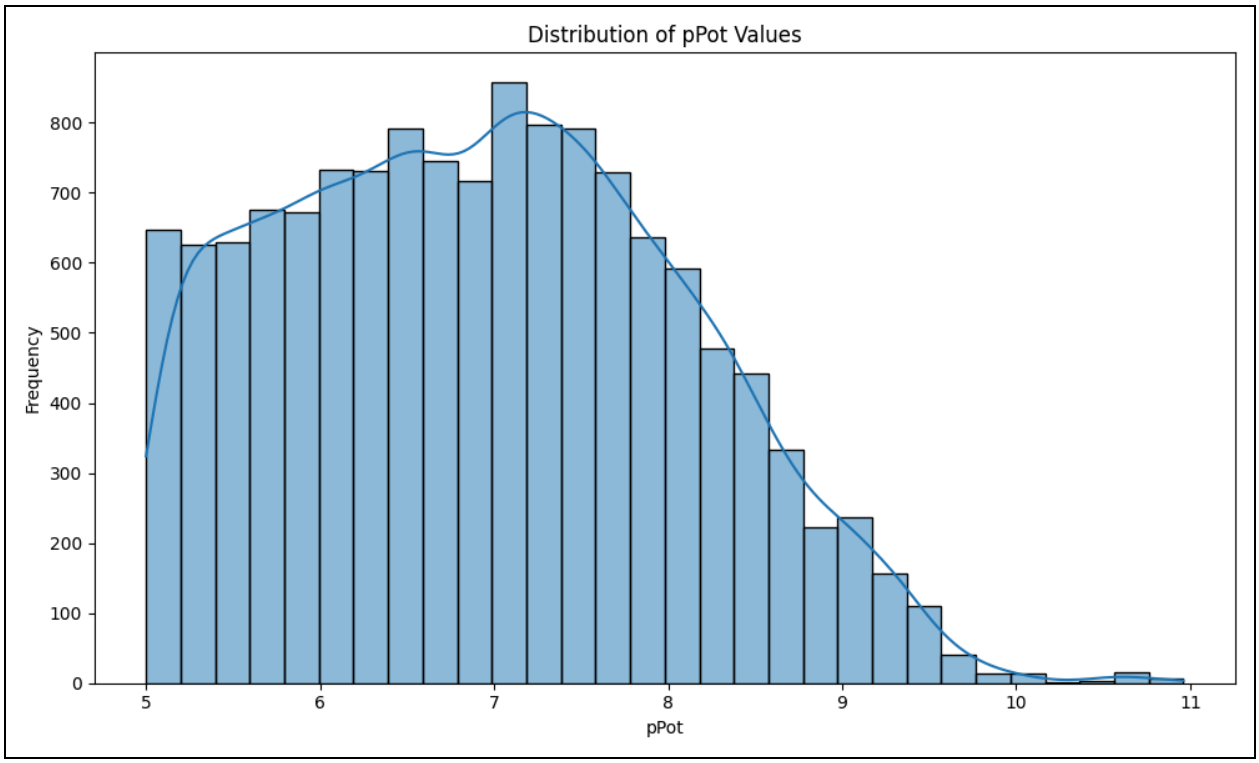
Distribution of pPot Values
Frequency
pPot
0
100
200
300
400
500
600
700
800
5
6
7
8
9
10
11

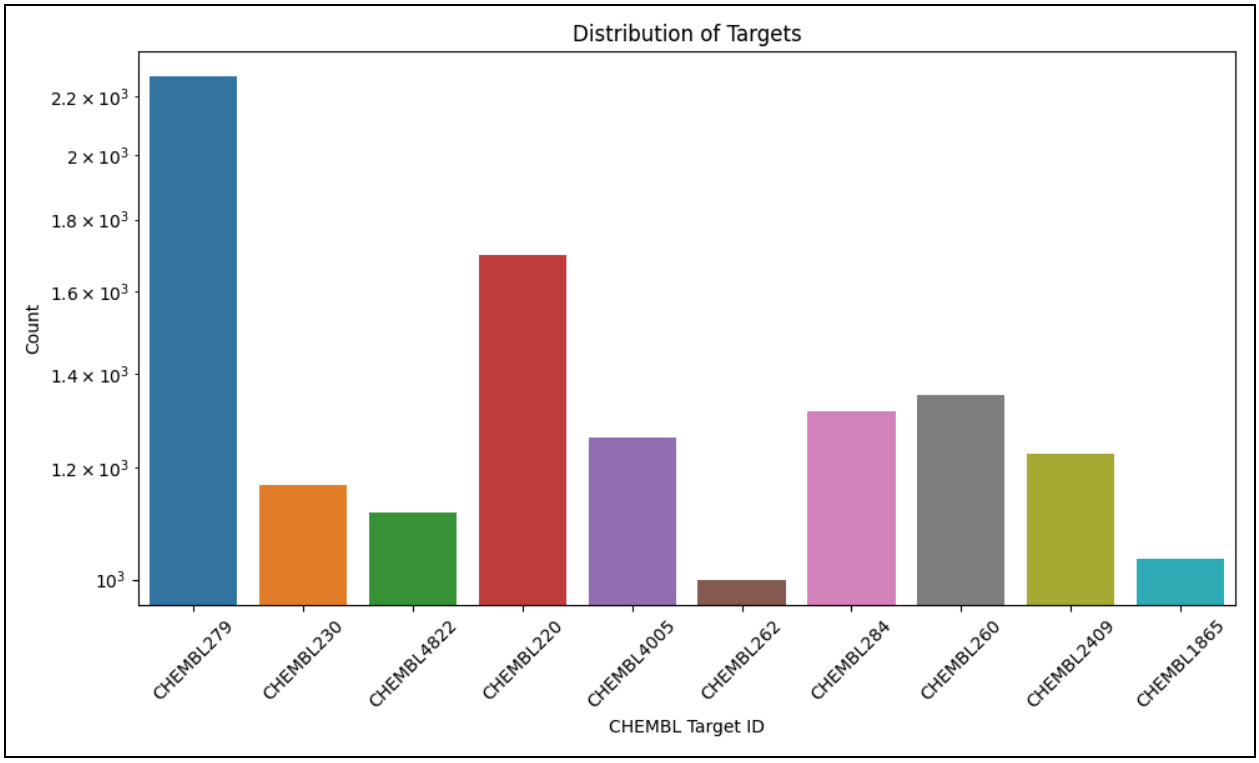
Distribution of Targets
Count
CHEMBL Target ID
2.2×10^3
2×10^3
1.8×10^3
1.6×10^3
1.4×10^3
1.2×10^3
10^3
CHEMBL279
CHEMBL230
CHEMBL4822
CHEMBL220
CHEMBL4005
CHEMBL262
CHEMBL284
CHEMBL260
CHEMBL2409
CHEMBL1865

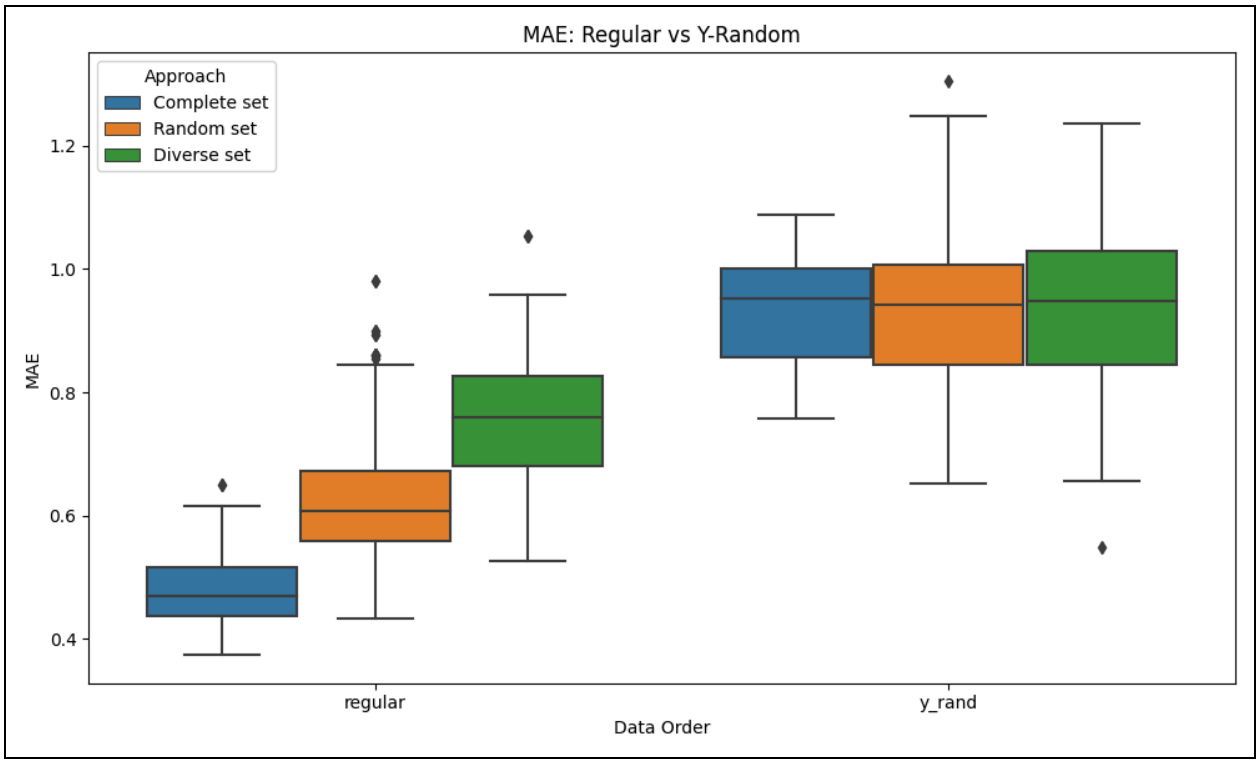

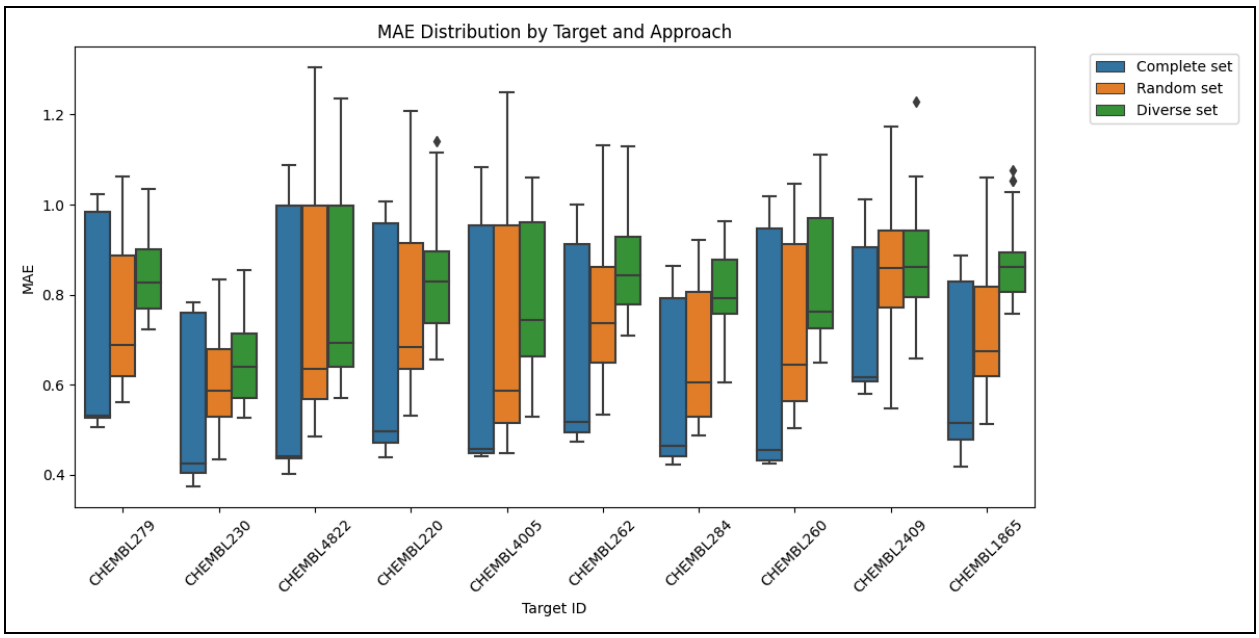

10. Future Directions

The team proposed several innovative directions to advance virtual screening:

Methodological Innovations

- **Hybrid AI/Physics**: Integrate deep learning with physics-based tools like RosettaVS for more accurate binding predictions.
- **Multi-Omics**: Combine genomics, proteomics, and metabolomics data to enhance prediction specificity.
- **Quantum Mechanics**: Leverage density functional theory (DFT)-based quantum computing for precise molecular modeling.
- **Active Learning**: Use ActiveDelta to prioritize compounds likely to be high-potency hits, optimizing screening efficiency.

Enhanced Benchmarking

- **Metrics**: Shift focus to top-10% hit rates (identifying the best compounds) rather than MAE or RMSE.
- **Testing**: Use distinct subsets (e.g., analog series) to reduce bias in model evaluation.

Applications

- **Drug Discovery**: Target “undruggable” proteins like KLHDC2 and NaV1.7, which are challenging but high-impact.
- **Personalized Medicine**: Tailor predictions using pharmacogenomics for patient-specific treatments.
- **Sustainable Chemistry**: Screen eco-friendly compounds to support green chemistry initiatives.

Broader Impact

- **Cost Reduction**: AI-driven virtual screening could reduce drug discovery costs by **40–60%**.
- **Global Health**: Accelerate antiviral development to address urgent health crises.

11. Conclusion

Key Takeaways

- **Original Study Validated**: Simple kNN models matched the accuracy of complex SVR and RFR models (MAE: 0.7–1.2 log units), exposing limitations in current benchmarking practices.
- **Replication Success**: The team improved MAE scores for several targets (e.g., 0.49 vs. 0.56 for VEGF receptor) using SPFP and deterministic training.
- **Practical Advancements**: Innovations like the Delta Classifier and meta-learning transformers increased top-tier hit rates by **16%** in low-data scenarios.

Significance

- **Rigorous Benchmarking**: Emphasized the need for control models (kNN, randomized predictions) to accurately assess machine learning performance.
- **Translational Potential**: Enhanced models reduce the need for extensive experimental validation, accelerating lead optimization in drug discovery.

Final Statement

By integrating hybrid AI methods, robust benchmarking, and multi-omics data, our advancements pave the way for faster, more cost-effective discovery of high-potency therapeutics, ushering in a new era for computational drug design.

11. References

1. Janela, T., Bajorath, J. (2022). Simple nearest-neighbour analysis meets the accuracy of compound potency predictions using complex machine learning models. *Nature Machine Intelligence*, 4, 1246–1255. https://doi.org/10.1038/s42256-022-00581-6
2. Additional citation: https://ouci.dntb.gov.ua/en/works/4zeGQOE7/
3. GitHub Repository: https://github.com/TiagoJanela/ML-for-compound-potency-prediction